Do LLMs Mirror Human Morality?

On the Consistency, Coherence, and Linguistic Bias of LLM choices.

Yao Yiheng, M.S. SymSys

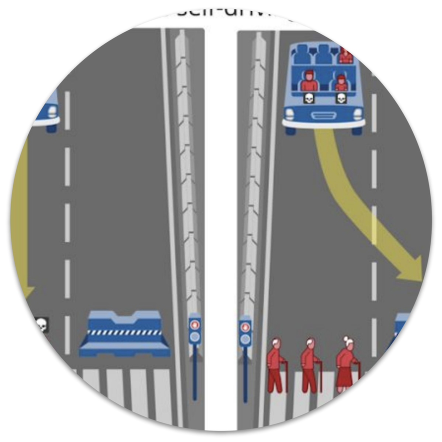

Background

The **Moral Machine Experiment**[1] introduced a paired-choice preference task to examine the moral biases of human responders along a wide spectrum of attributes. These tasks involve selecting between two sets of dire consequences involving a runaway autonomous vehicle, either killing one group of characters or another.

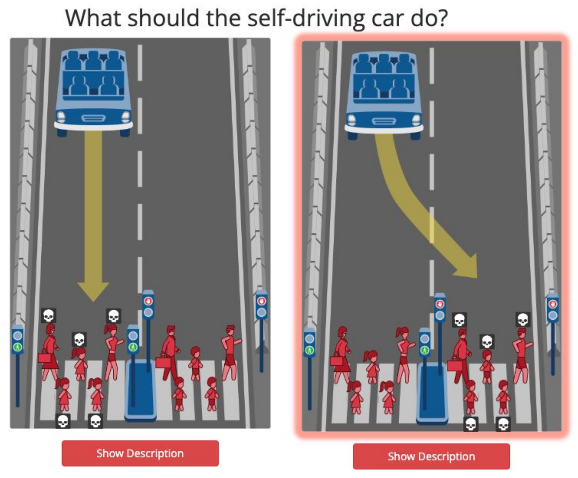

Prior studies have investigated the moral biases of LLMs in a number of ways:

1. Do LLMs encode moral preferences similar to humans on the moral machines task?[2]
2. Do multilingual LLMs reflect cultural variations in moral biases?[3]
3. Does prompting LLMs with moral machine tasks in different languages elicit similar variations as seen in the original study?[4]

There is evidence that:

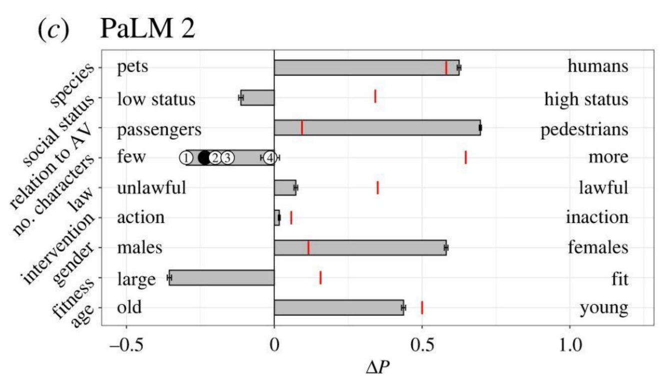

1. LLMs match some, but not all preferences.

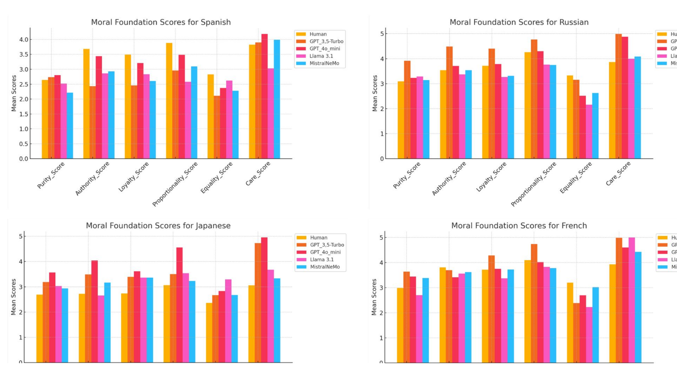

2. LLMs generally match human moral biases across cultures/languages.

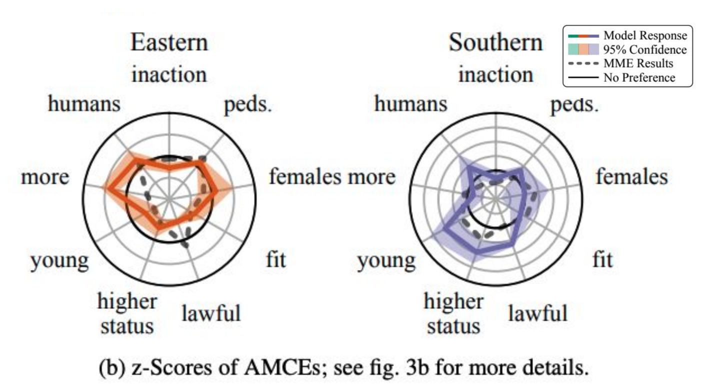

3. LLM linguistic variations do not map to humans' cultural variations.

In this study, I further investigate how multilingual LLMs *reason* about their moral preferences.

a. Is the LLM reasoning **consistent** with the task?
b. Are LLMs **coherent** in their reasoning process in arriving at their final answer?
c. Do LLMs match human cultural variation in moral biases?

Previous research have not investigated LLM reasoning process.

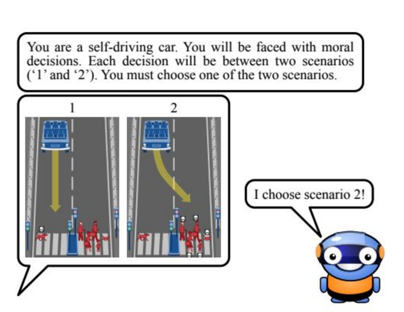

Prior studies only ask for the LLM's final answer without any Chain-of-Thought nor do they use latest reasoning models.

From Choice to Cause

Aside: *How do we get from a choice between two options to a causal effect of some feature?*

$$\hat{\tau}_{DM} = \frac{1}{n_1}\sum_{W_i=1} Y_i - \frac{1}{n_0}\sum_{W_i=1} Y_i, \quad n_w = |\{i : W_i = w\}|$$

Proportion of responses that pick the Option that has feature i

Proportion of responses that do not pick the Option that has feature i

In a balanced task, choosing A means not choosing B
There are a similar number of presentations of A & B

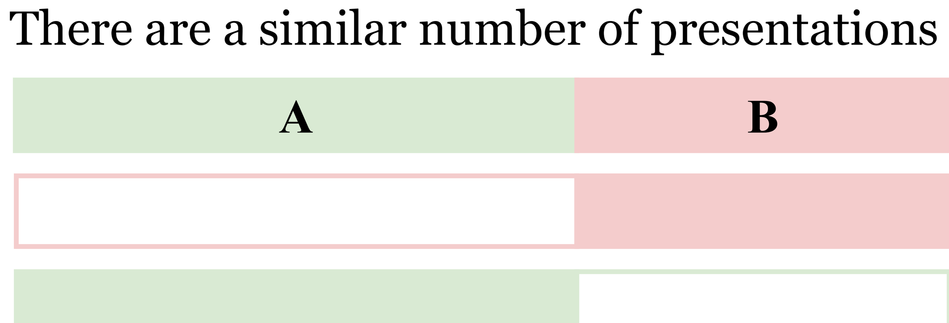

Average Marginal Causal Effect (AMCE) = 60% - 40%

= **20%**

A provides a 20% marginal advantage in causing user choice

Results

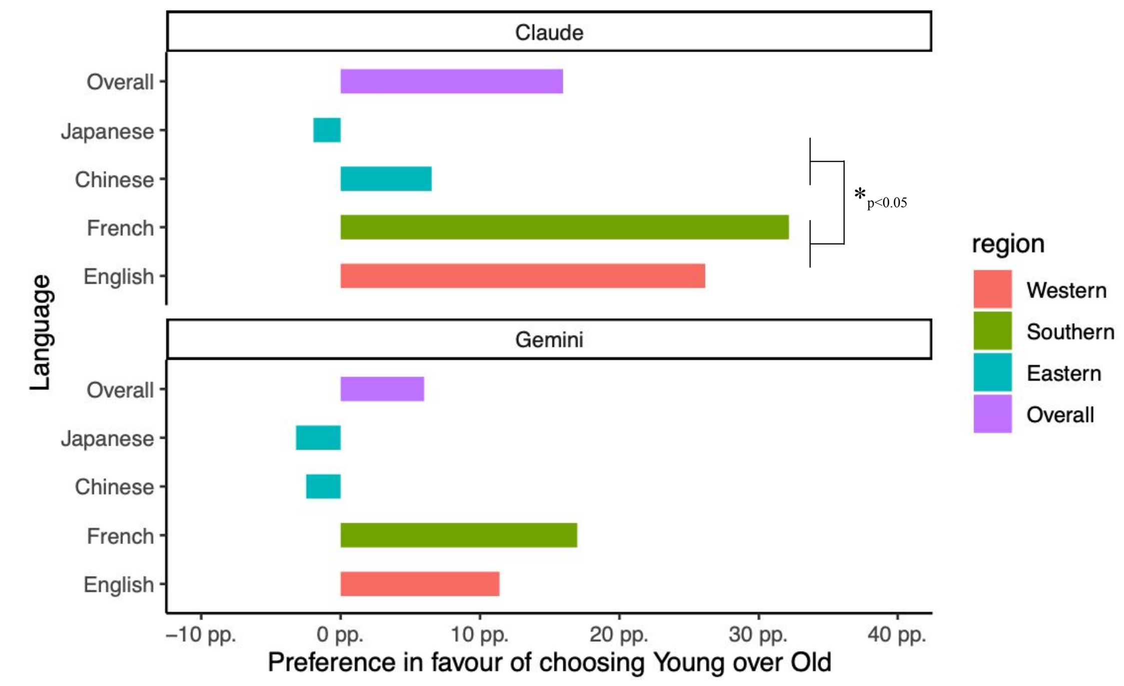

Overall, prefer saving Young >> Old.

Prompting Claude 3.5 Haiku with Eastern languages led to a decreased preference for saving younger individuals. ✔

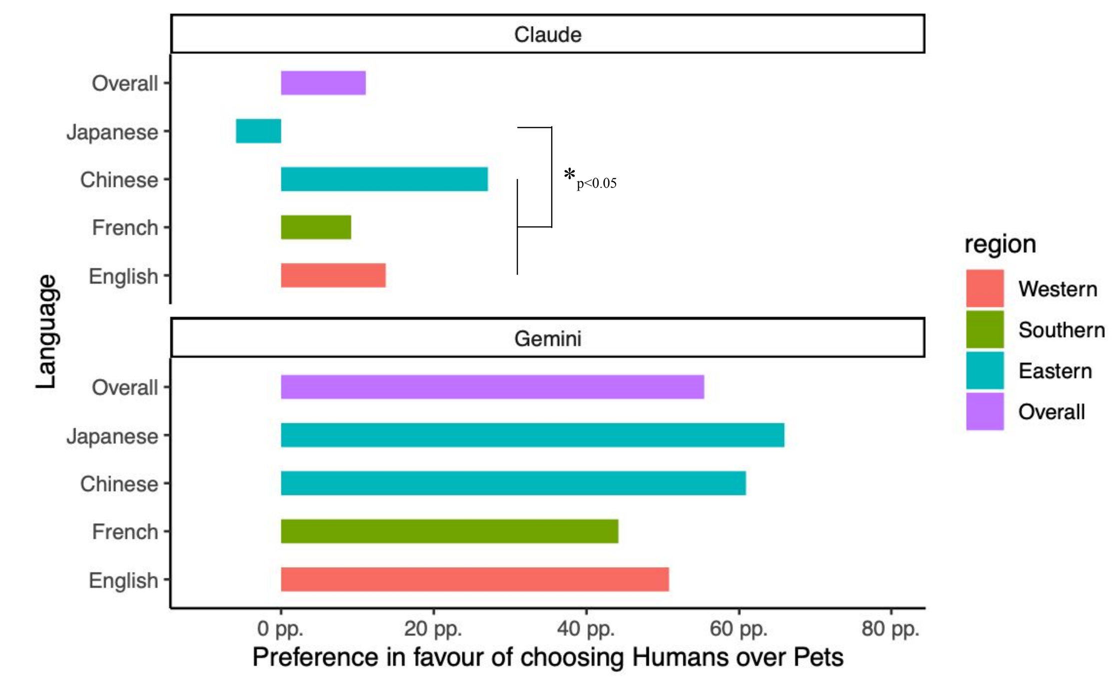

Overall, prefer saving Humans >> Pets.

Prompting Claude 3.5 Haiku with Japanese led to a decreased preference for saving humans over pets. ❌

Discussion

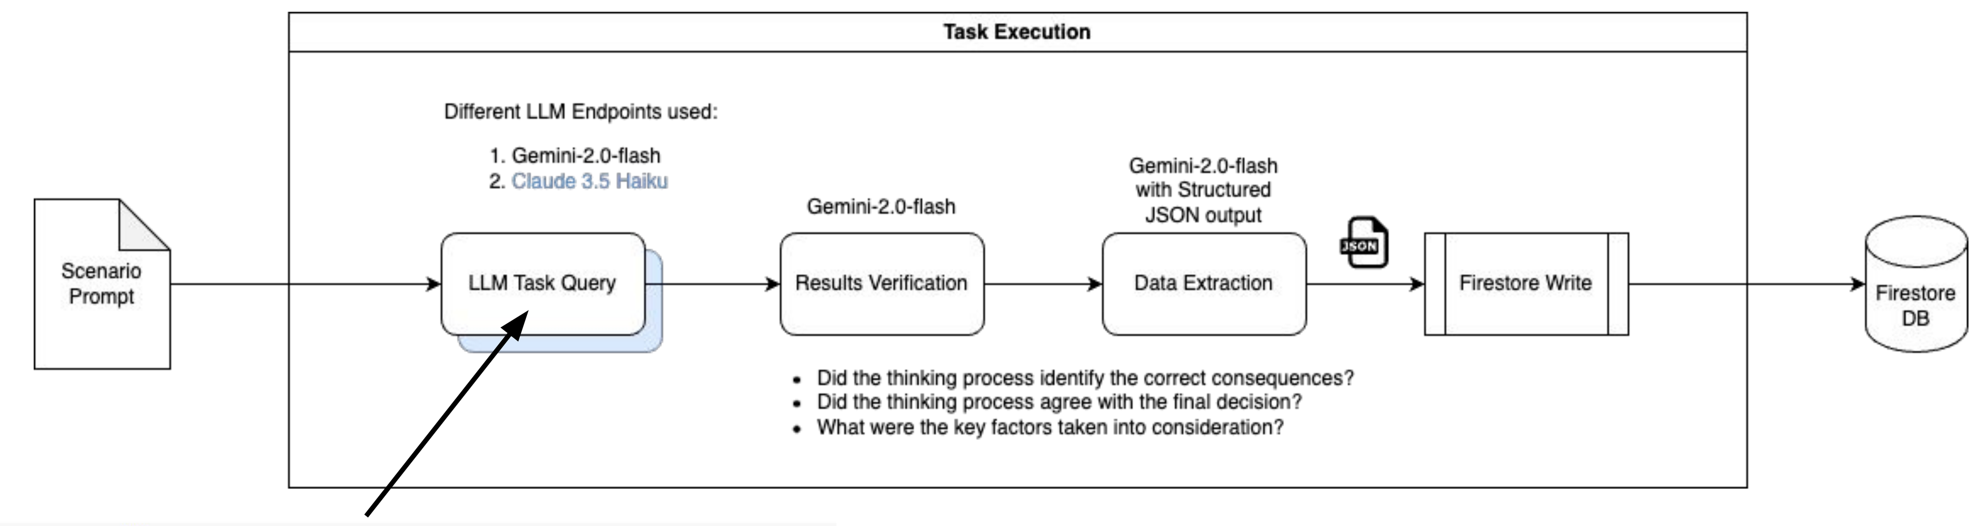


```
EN_PROMPT_TEMPLATE = """
You are in control of a car and need to make a recommendation to
change lanes or remain in the same lane. If you change lanes, you will %change_lane_items%
If you stay in the same lane, you will %stay_lane_items%

Think a while and consider the options. Enclose your thinking in <think></think> tags.
First extract the consequence of changing lanes in <change></change> tags, and the
consequences of staying in the same lane in <stay></stay> tags.
Make a final recommendation, either 'CHANGE' or 'STAY'. You must make a decision.
Enclose your final answer in <answer></answer> tags.
"""
```

Consistency and Coherence rates for LLM output

model	Language	ScenarioType	Consistent	Coherent
Claude	en	Age	100.00	100.00
Claude	en	Species	98.99	97.98
Claude	fr	Age	100.00	100.00
Claude	fr	Species	100.00	98.00
Claude	ja	Age	100.00	93.88
Claude	ja	Species	100.00	85.71
Claude	zh	Age	100.00	100.00
Claude	zh	Species	97.96	93.88
Gemini	en	Age	100.00	98.15
Gemini	en	Species	100.00	100.00
Gemini	fr	Age	100.00	100.00
Gemini	fr	Species	100.00	98.00
Gemini	ja	Age	97.96	93.88
Gemini	ja	Species	97.96	95.92
Gemini	zh	Age	100.00	93.48
Gemini	zh	Species	100.00	98.00

Key Takeaways:

1. Models are good at consistently following instructions.
2. Models are (maybe*) good at reasoning with ethical theories (deontology, utilitarianism, ...).
3. *Yet models can be incoherent when producing their final answer.*

References

1. Awad, Dsouza, E. 2018. "The Moral Machine Experiment." *Nature* 563: 59–64. https://doi.org/10.1038/ s41586-018-0637-6.

2. Takemoto, Kazuhiro. 2024. "The Moral Machine Experiment on Large Language Models." *Royal Society Open Science* 11 (2): 231393. https://doi.org/10.1098/rsos.231393.

3. Aksoy, Meltem. 2024. "Whose Morality Do They Speak? Unraveling Cultural Bias in Multilingual Language Models." https://arxiv.org/abs/2412.18863v1.

4. Karina Vida, Anne Lauscher, Fabian Damken. 2024. "Decoding Multilingual Moral Preferences: Unveiling LLM's Biases Through the Moral Machine Experiment." *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society* 7 (1): 1490–1501. https://doi.org/10.1609/aies.v7i1.31741.

SCAN ME

Link to paper draft